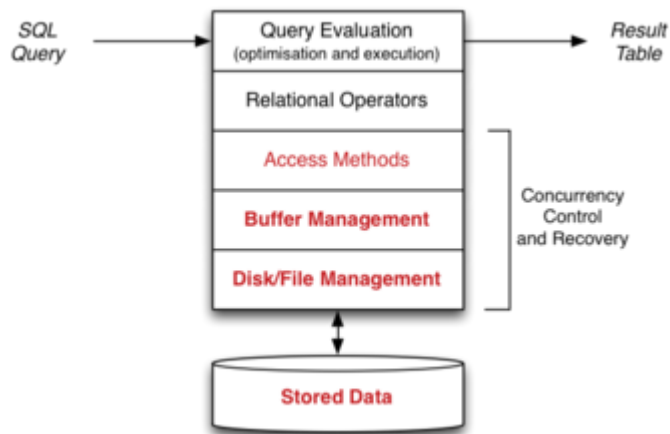# Week 02 Lectures

Storage Manager

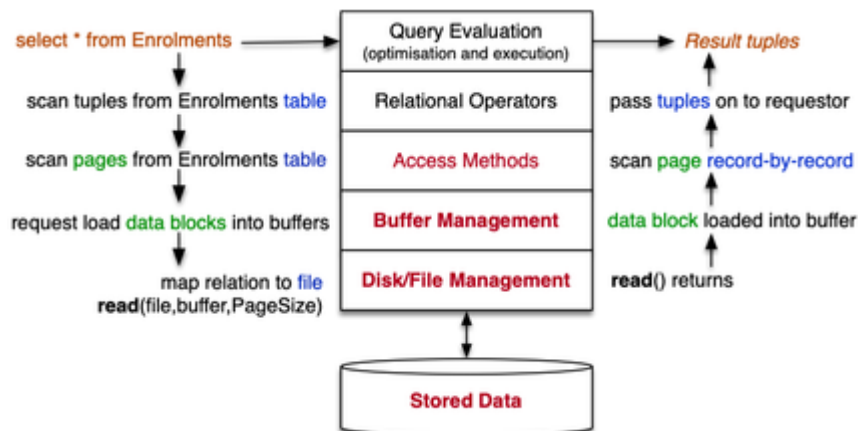### Storage Management

Levels of DBMS related to storage management:

Aims of storage management in DBMS:

- map from database objects (e.g. tables) to disk files
- manage transfer of data to/from disk storage
- use buffers to minimise disk/memory transfers
- interpret loaded data as tuples/records
- provide view of data as collection of pages/tuples
- basis for "file structures" used by access methods

Views of Data in Query Evaluation

Representing database objects during query execution:

- `DB` (handle on an authorised/opened database)
- `Rel` (handle on an opened relation)
- `Page` (memory buffer to hold contents of disk block)
- `Tuple` (memory holding data values from one tuple)

Addressing in DBMSs:

- `PageID = FileID+Offset` ... identifies a block of data
    - where `Offset` gives location of block within file
- `TupleID = PageID+Index` ... identifies a single tuple
    - where `Index` gives access to location of tuple within page

---

Storage Management

Topics in storage management ...

- Disks and Files
    - performance issues and organisation of disk files
- Buffer Management
    - using caching to improve DBMS system throughput
- Tuple/Page Management
    - how tuples are represented within disk pages
- DB Object Management (Catalog)
    - how tables/views/functions/types, etc. are represented

Storage Technology

Storage Technology

Persistent storage is

- large, cheap, relatively slow, accessed in blocks
- used for long–term *storage* of data

Computational storage is

- small, expensive, fast, accessed by byte/word
- used for all *analysis*/*calculation* of data

Access cost HDD:RAM $\cong$ 100000:1, e.g.

- 10ms to read block containing two tuples
- 1μs to compare fields in two tuples

... Storage Technology

Hard disks are well–established, cheap, high–volume, ...

Alternative bulk storage: SSD

- faster than HDDs, no latency
- can read single items
- update requires block erase then write
- over time, writes "wear out" blocks

- require controllers that spread write load

Feasible for long–term, high–update environments?

Comparison of HDD and SSD properties:

	HDD	SDD
Cost/byte	~ 2c / GB	~ 13c / GB
Read latency	~ 10ms	~ 50µs
Write latency	~ 10ms	~ 900µs
Read unit	block (e.g. 1KB)	byte
Writing	write a block	write on empty block

Will SSDs ever replace HDDs for large–scale database storage?

Cost Models

Throughout this course, we compare costs of DB operations

Important aspects in determining cost:

- data is always transferred to/from disk as whole blocks (pages)
- cost of manipulating tuples in memory is negligible
- overall cost determined primarily by #data–blocks read/written

Complicating factors in determining costs:

- not all page accesses require disk access (buffer pool)
- tuples typically have variable size (tuples/page ?)

More details later ...

File Management

Aims of file management subsystem:

- organise layout of data within the filesystem
- handle mapping from database ID to file address
- transfer blocks of data between buffer pool and filesystem
- also attempts to handle file access error problems (retry)

Builds higher–level operations on top of OS file operations.

Typical file operations provided by the operating system:

```
fd = open(fileName,mode)
  // open a named file for reading/writing/appending
close(fd)
  // close an open file, via its descriptor
nread = read(fd, buf, nbytes)
  // attempt to read data from file into buffer
nwritten = write(fd, buf, nbytes)
  // attempt to write data from buffer to file
lseek(fd, offset, seek_type)
  // move file pointer to relative/absolute file offset
fsync(fd)
  // flush contents of file buffers to disk
```

DBMS File Organisation

How is data for DB objects arranged in the file system?

Different DBMSs make different choices, e.g.

- by–pass the file system and use a raw disk partition
- have a single very large file containing all DB data
- have several large files, with tables spread across them
- have multiple data files, one for each table
- have multiple files for each table
- etc.

Single–file Storage Manager

Consider the following simple single–file DBMS layout:

E.g.

SpaceMap = [(0,10,U), (10,10,U), (20,600,U), (620,100,U), (720,20,F)]

NameMap = [("employee",20,350), ("project",620,40)]

... Single–file Storage Manager

Each file segment consists of a number fixed–size blocks

The following data/constant definitions are useful

```
#define PAGESIZE 4096   // bytes per page

typedef long PageId;    // PageId is block index
                        // pageOffset=PageId*PAGESIZE

typedef char *Page;     // pointer to page/block buffer
```

Typical PAGESIZE values: 1024, 2048, 4096, 8192

... Single–file Storage Manager

Storage Manager data structures for opened DBs & Tables

```c
typedef struct DBrec {
    char *dbname;    // copy of database name
    int fd;          // the database file
    SpaceMap map;    // map of free/used areas
    NameMap names;   // map names to areas + sizes
} *DB;

typedef struct Relrec {
    char *relname;   // copy of table name
    int   start;     // page index of start of table data
    int   npages;    // number of pages of table data
    ...
} *Rel;
```

Example: Scanning a Relation

With the above disk manager, our example:

```
select name from Employee
```

might be implemented as something like

```c
DB db = openDatabase("myDB");
Rel r = openRelation(db,"Employee");
Page buffer = malloc(PAGESIZE*sizeof(char));
for (int i = 0; i < r->npages; i++) {
    PageId pid = r->start+i;
    get_page(db, pid, buffer);
    for each tuple in buffer {
        get tuple data and extract name
        add (name) to result tuples
    }
}
```

Single–File Storage Manager

```c
// start using DB, buffer meta-data
DB openDatabase(char *name) {
    DB db = new(struct DBrec);
    db->dbname = strdup(name);
    db->fd = open(name,O_RDWR);
    db->map = readSpaceMap(db->fd);
    db->names = readNameMap(db->fd);
    return db;
}
// set up struct describing relation
Rel openRelation(DB db, char *rname) {
    Rel r = new(struct Relrec);
    r->relname = strdup(rname);
    // get relation data from map tables
    r->start = ...;
    r->npages = ...;
    return r;
}
```

... Single–File Storage Manager

```c
// assume that Page = byte[PageSize]
// assume that PageId = block number in file

// read page from file into memory buffer
void get_page(DB db, PageId p, Page buf) {
    lseek(db->fd, p*PAGESIZE, SEEK_SET);
    read(db->fd, buf, PAGESIZE);
}
```

```
// write page from memory buffer to file
void put_page(Db db, PageId p, Page buf) {
    lseek(db->fd, p*PAGESIZE, SEEK_SET);
    write(db->fd, buf, PAGESIZE);
}
```

Exercise 1: Relation Scan Cost

Consider a table $R(x,y,z)$ with 10^5 tuples, implemented as

- number of tuples $r = 10,000$
- average size of tuples $R = 200$ bytes
- size of data pages $B = 4096$ bytes
- time to read one data page $T_r = 10msec$

- time to check one tuple 1 usec
- time to form one result tuple 1 usec
- time to write one result page $T_r = 10msec$

Calculate the total time–cost for answering the query:

```
insert into S select * from R where x > 10;
```

if 50% of the tuples satisfy the condition.

DBMS Parameters

Our view of relations in DBMSs:

- a relation is a set of r tuples, with average size R bytes
- the tuples are stored in b data pages on disk
- each page has size B bytes and contains up to c tuples
- data is transferred disk↔memory in whole pages
- cost of disk↔memory transfer T_r, T_w dominates other costs

... DBMS Parameters

Typical DBMS/table parameter values:

Quantity	Symbol	E.g. Value
total # tuples	r	10^6
record size	R	128 bytes
total # pages	b	10^5
page size	B	8192 bytes
# tuples per page	c	60
page read/write time	T_r, T_w	10 msec
cost to process one page in memory	–	$\cong 0$

Multiple–file Disk Manager

Most DBMSs don't use a single large file for all data.

They typically provide:

- multiple files partitioned physically or logically
- mapping from DB–level objects to files (e.g. via meta–data)

Precise file structure varies between individual DBMSs.

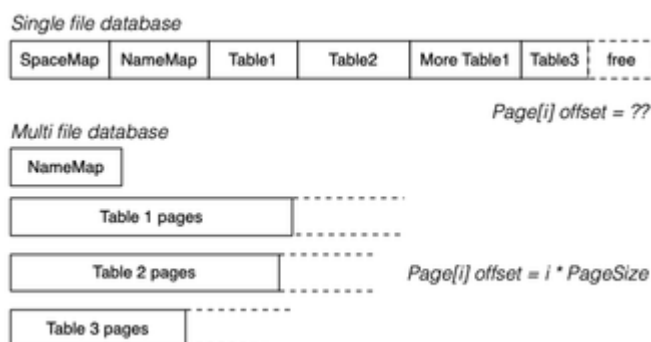Using multiple files (one file per relation) can be easier, e.g.

- adding a new relation
- extending the size of a relation
- computing page offsets within a relation

Example of single–file vs multiple–file:

Consider how you would compute file offset of page[i] in table[1] ...

Structure of `PageId` for data pages in such systems ...

If system uses one file per table, `PageId` contains:

- relation identifier (which can be mapped to filename)
- page number (to identify page within the file)

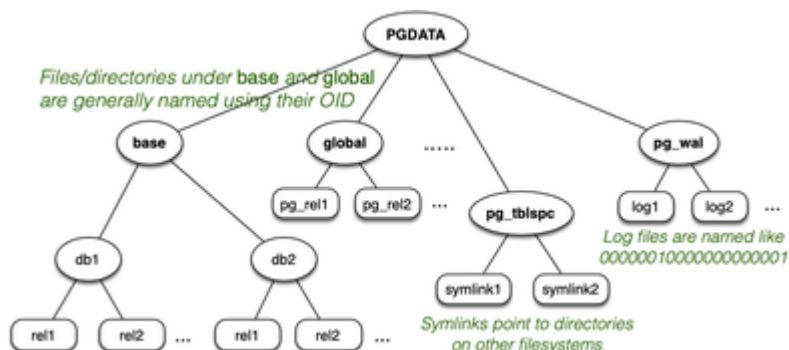If system uses several files per table, `PageId` contains:

- relation identifier
- file identifier (combined with relid, gives filename)
- page number (to identify page within the file)

PostgreSQL Storage Manager

PostgreSQL uses the following file organisation ...

Components of storage subsystem:

- mapping from relations to files (**RelFileNode**)
- abstraction for open relation pool (**storage/smgr**)
- functions for managing files (**storage/smgr/md.c**)
- file–descriptor pool (**storage/file**)

PostgreSQL has two basic kinds of files:

- heap files containing data (tuples)
- index files containing index entries

Note: `smgr` designed for many storage devices; only disk handler provided

---

Relations as Files

PostgreSQL identifies relation files via their OIDs.

The core data structure for this is **RelFileNode**:

```
typedef struct RelFileNode {
    Oid  spcNode;  // tablespace
    Oid  dbNode;   // database
    Oid  relNode;  // relation
} RelFileNode;
```

Global (shared) tables (e.g. `pg_database`) have

- spcNode == GLOBALTABLESPACE_OID
- dbNode == 0

... Relations as Files

The **relpath** function maps **RelFileNode** to file:

```
char *relpath(RelFileNode r)  // simplified
{
   char *path = malloc(ENOUGH_SPACE);

   if (r.spcNode == GLOBALTABLESPACE_OID) {
      /* Shared system relations live in PGDATA/global */
      Assert(r.dbNode == 0);
      sprintf(path, "%s/global/%u",
              DataDir, r.relNode);
   }
   else if (r.spcNode == DEFAULTTABLESPACE_OID) {
      /* The default tablespace is PGDATA/base */
      sprintf(path, "%s/base/%u/%u",
              DataDir, r.dbNode, r.relNode);
   }
   else {
      /* All other tablespaces accessed via symlinks */
      sprintf(path, "%s/pg_tblspc/%u/%u/%u", DataDir
              r.spcNode, r.dbNode, r.relNode);
   }
   return path;
}
```

Exercise 2: PostgreSQL Files

In your PostgreSQL server

- examine the content of the `$PGDATA` directory
- find the directory containing the `pizza` database
- find the file in this directory for the `People` table

- examine the contents of the `People` file
- what are the other files in the directory?
- are there *forks* in any of your databases?

File Descriptor Pool

Unix has limits on the number of concurrently open files.

PostgreSQL maintains a pool of open file descriptors:

- to hide this limitation from higher level functions
- to minimise expensive `open()` operations

File names are simply strings: **`typedef char *FileName`**

Open files are referenced via: **`typedef int File`**

A **`File`** is an index into a table of "virtual file descriptors".

Defs: **`include/storage/fd.h`**
Code: **`backend/storage/file/fd.c`**

... File Descriptor Pool

Interface to file descriptor (pool):

```
File PathNameOpenFile(char *fileName, int flags)
     // open a file with default pg.conf mode
File PathNameOpenFilePerm(char *fName, int flags, int mode)
     // open a file in the DB directory ($PGDATA/base/...)
File OpenTemporaryFile(bool interXact)
     // open temp file flag: close at end of transaction?
void FileClose(File file)
void FileUnlink(File file)
int  FileRead(File file, char *buffer, int amount)
int  FileWrite(File file, char *buffer, int amount)
int  FileSync(File file)
long FileSeek(File file, long offset, int whence)
int  FileTruncate(File file, long offset)
```

Analogous to Unix syscalls `open()`, `close()`, `read()`, ...

... File Descriptor Pool

Virtual file descriptor records (simplified):

```
typedef struct vfd
{
    s_short  fd;               // current FD, or VFD_CLOSED if none
    u_short  fdstate;          // bitflags for Vfd's state
    File     nextFree;         // link to next free Vfd, if in freelist
    File     lruMoreRecently;  // doubly linked recency-of-use list
    File     lruLessRecently;
    long     seekPos;          // current logical file position
    char     *fileName;        // name of file, or NULL for unused Vfd
    // NB: fileName is malloc'd, and must be free'd when closing the Vfd
    int      fileFlags;        // open(2) flags for (re)opening the file
    int      fileMode;         // mode to pass to open(2)
} Vfd;
```

... File Descriptor Pool

Virtual file descriptors (**`vfd`**)

- physically stored in dynamically-allocated array

- also arranged into list by recency–of–use

`VfdCache[0]` holds list head/tail pointers.

Exercise 3: Opening a Vfd

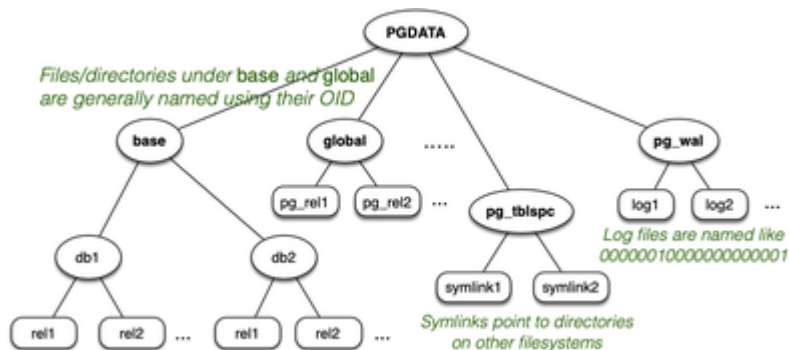Consider the following call to open a file

```
f = PathNameOpenFilePerm(
    "/srvr/jas/pgsql/data/base/13645/12348",
    O_RDWR | O_CREAT | O_EXCL | PG_BINARY,
    0600
)
```

Sketch implementation of `PathNameOpenFilePerm()`

File Manager

Reminder: PostgreSQL file organisation
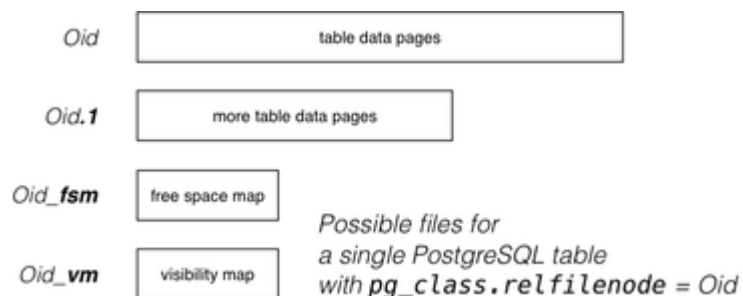

... File Manager

PostgreSQL stores each table

- in the directory *PGDATA*/pg_database.oid
- often in multiple files (aka *forks*)


... File Manager

Data files (*Oid*, *Oid*.1, ...):

- sequence of fixed–size blocks/pages (typically 8KB)
- each page contains tuple data and admin data (see later)
- max size of data files 1GB (Unix limitation)

PostgreSQL Data File (Heap)

Free space map   (*Oid*_fsm):

- indicates where free space is in data pages
- "free" space is only free after VACUUM
    (DELETE simply marks tuples as no longer in use xmax)

Visibility map   (*Oid*_vm):

- indicates pages where all tuples are "visible"
    (*visible* = accessible to all currently active transactions)
- such pages can be ignored by VACUUM

The "magnetic disk storage manager" (**storage/smgr/md.c**)

- manages its own pool of open file descriptors (Vfd's)
- may use several Vfd's to access data, if several forks
- manages mapping from **PageID** to file+offset.

PostgreSQL PageID values are structured:

```
typedef struct
{
    RelFileNode rnode;     // which relation/file
    ForkNumber  forkNum;   // which fork (of reln)
    BlockNumber blockNum; // which page/block
} BufferTag;
```

Access to a block of data proceeds (roughly) as follows:

```
// pageID set from pg_catalog tables
// buffer obtained from Buffer pool
getBlock(BufferTag pageID, Buffer buf)
{
   File fid;  off_t offset;  int fd;
   (fid, offset) = findBlock(pageID)
   fd = VfdCache[fid].fd;
   lseek(fd, offset, SEEK_SET)
   VfdCache[fid].seekPos = offset;
   nread = read(fd, buf, BLOCKSIZE)
   if (nread < BLOCKSIZE) ... we have a problem
}
```

BLOCKSIZE is a global configurable constant (default: 8192)

```
findBlock(BufferTag pageID) returns (Vfd, off_t)
{
   offset = pageID.blockNum * BLOCKSIZE
   fileName = relpath(pageID.rnode)
```

```
    if (pageID.forkNum > 0)
        fileName = fileName+"."+pageID.forkNum
    fid = PathNameOpenFIle(fileName, O_READ);
    fSize = VfdCache[fid].fileSize;
    if (offset > fSize) {
        fid = allocate new Vfd for next fork
        offset = offset - fd.fileSize
    }
    return (fd, offset)
}
```

Buffer Pool

Buffer Pool

Aim of buffer pool:

- hold pages read from database files, for possible re−use

Used by:

- *access methods* which read/write data pages
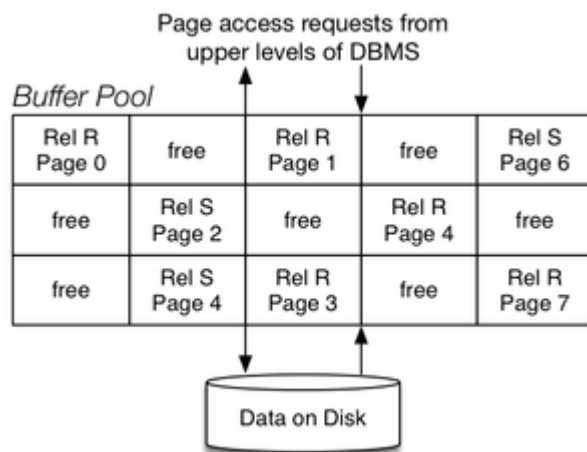- e.g. sequential scan, indexed retrieval, hashing

Uses:

- file manager functions to access data files

Note: we use the terms *page* and *block* interchangably

... Buffer Pool

... Buffer Pool

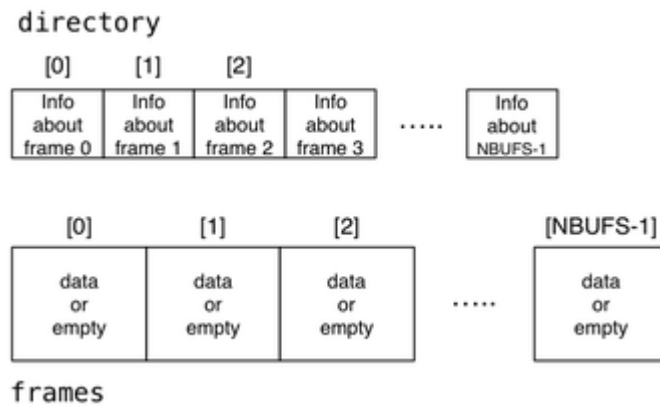Buffer pool operations: (both take single `PageID` argument)

- `request_page(pid)`, `release_page(pid)`, ...

To some extent ...

- `request_page()` replaces `getBlock()` or `get_page()`
- `release_page()` replaces `putBlock()` or `put_page()`

Buffer pool data structures:

- frames ... array of `NBUFS` Page buffers
- directory ... array of `NBUFS` FrameData items

For each frame, we need to know: (`FrameData`)

- which Page it contains, or whether empty/free
- whether it has been modified since loading (*dirty bit*)
- how many transactions are currently using it (*pin count*)
- time–stamp for most recent access (assists with replacement)

Pages are referenced by PageID ...

- PageID = BufferTag = (rnode, forkNum, blockNum)

How scans are performed without Buffer Pool:

```
Buffer buf;
int N = numberOfBlocks(Rel);
for (i = 0; i < N; i++) {
   pageID = makePageID(db,Rel,i);
   getBlock(pageID, buf);
   for (j = 0; j < nTuples(buf); j++)
      process(buf, j)
}
```

Requires `N` page reads.

If we read it again, `N` page reads.

How scans are performed with Buffer Pool:

```
Buffer buf;
int N = numberOfBlocks(Rel);
for (i = 0; i < N; i++) {
   pageID = makePageID(db,Rel,i);
   bufID = request_page(pageID);
   buf = frames[bufID]
   for (j = 0; j < nTuples(buf); j++)
      process(buf, j)
   release_page(pageID);
}
```

Requires `N` page reads on the first pass.

If we read it again, $0 \leq$ page reads $\leq$ `N`

Buffer pool data structures:

```
typedef char Page[PAGESIZE];
typedef ... PageID;  // defined earlier

typedef struct _FrameData {
   PageID  pid;         // which page is in frame
   int     pin_count;   // how many processes using page
   int     dirty;       // page modified since loaded?
   Time    last_used;   // when page was last accessed
} FrameData;

Page frames[NBUFS];      // actual buffers
FrameData directory[NBUFS];
```

... Buffer Pool

Implementation of `request_page()`

```
int request_page(PageID pid)
{
   bufID = findInPool(pid)
   if (pid == NOT_FOUND) {
      if (no free frames in Pool) {
         bufID = findFrameToReplace()
         if (directory[bufID].dirty)
            old = directory[bufID].page
            put_page(old, frames[bufID])
      }
      bufID = index of freed frame
      directory[bufID].page = pid
      directory[bufID].pin_count = 0
      directory[bufID].dirty = 0
      get_page(pid, frames[bufID])
   }
   directory[bufID].pin_count++
   return bufID
}
```

... Buffer Pool

The `release_page(pid)` operation:

- Decrement pin count for specified page

Note: no effect on disk or buffer contents until replacement required

The `mark_page(pid)` operation:

- Set dirty bit on for specified page

Note: doesn't actually write to disk; indicates that page changed

The `flush_page(pid)` operation:

- Write the specified page to disk (using `write()`)

Note: not generally used by higher levels of DBMS

... Buffer Pool

Evicting a page ...

- find frame(s) preferably satisfying
  - pin count = 0 (i.e. nobody using it)
  - dirty bit = 0 (not modified)
- if selected frame was modified, flush frame to disk
- flag directory entry as "frame empty"

If multiple frames can potentially be released

- need a policy to decide which is best choice

Page Replacement Policies

Several schemes are commonly in use:

- Least Recently Used (LRU)
- Most Recently Used (MRU)
- First in First Out (FIFO)
- Random

LRU / MRU require knowledge of when pages were last accessed

- how to keep track of "last access" time?
- base on request/release ops or on *real* page usage?

... Page Replacement Policies

Cost benefit from buffer pool (with n frames) is determined by:

- number of available frames (more $\Rightarrow$ better)
- replacement strategy vs page access pattern

Example (a): sequential scan, LRU or MRU, $n \geqslant b$

First scan costs b reads; subsequent scans are "free".

Example (b): sequential scan, MRU, $n < b$

First scan costs b reads; subsequent scans cost $b - n$ reads.

Example (c): sequential scan, LRU, $n < b$

All scans cost b reads; known as *sequential flooding*.

Effect of Buffer Management

Consider a query to find customers who are also employees:

```
select c.name
from    Customer c, Employee e
where   c.ssn = e.ssn;
```

This might be implemented inside the DBMS via nested loops:

```
for each tuple t1 in Customer {
    for each tuple t2 in Employee {
        if (t1.ssn == t2.ssn)
            append (t1.name) to result set
    }
}
```

... Effect of Buffer Management

In terms of page–level operations, the algorithm looks like:

```
Rel rC = openRelation("Customer");
Rel rE = openRelation("Employee");
for (int i = 0; i < nPages(rC); i++) {
    PageID pid1 = makePageID(db,rC,i);
    Page p1 = request_page(pid1);
    for (int j = 0; j < nPages(rE); j++) {
        PageID pid2 = makePageID(db,rE,j);
        Page p2 = request_page(pid2);
        // compare all pairs of tuples from p1,p2
        // construct solution set from matching pairs
        release_page(pid2);
    }
    release_page(pid1);
}
```

Exercise 4: Buffer Cost Benefit (i)

Assume that:

- the `Customer` relation has b_C pages (e.g. 10)
- the `Employee` relation has b_E pages (e.g. 4)

Compute how many page reads occur ...

- if we have only 2 buffers (i.e. effectively no buffer pool)
- if we have 20 buffers
- when a buffer pool with MRU replacement strategy is used
- when a buffer pool with LRU replacement strategy is used

For the last two, buffer pool has $n=3$ slots ($n < b_C$ and $n < b_E$)

Exercise 5: Buffer Cost Benefit (ii)

If the tables were larger, the above analysis would be tedious.

Write a C program to simulate buffer pool usage

- assuming a nested loop join as above
- `argv[1]` gives number of pages in "outer" table
- `argv[2]` gives number of pages in "inner" table
- `argv[3]` gives number of slots in buffer pool
- `argv[4]` gives replacement strategy (LRU,MRU,FIFO–Q)

PostgreSQL Buffer Manager

PostgreSQL buffer manager:

- provides a shared pool of memory buffers for all backends
- all access methods get data from disk via buffer manager

Buffers are located in a large region of *shared memory*.

Definitions: **src/include/storage/buf*.h**

Functions: **src/backend/storage/buffer/*.c**

Commentary: **backend/storage/buffer/README**

Buffer code is also used by backends who want a private buffer pool

... PostgreSQL Buffer Manager

Buffer pool consists of:

BufferDescriptors

- shared fixed array (size `NBuffers`) of **BufferDesc**

BufferBlocks

- shared fixed array (size `NBuffers`) of **Buffer**

Buffer = index in above arrays

- indexes: global buffers `1..NBuffers`; local buffers negative

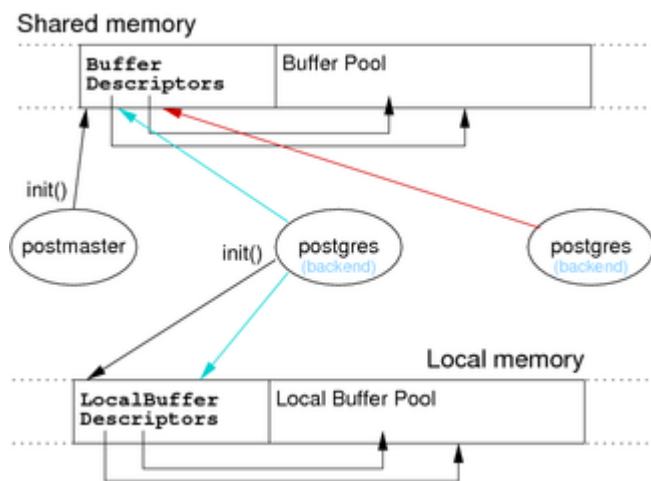Size of buffer pool is set in *postgresql.conf*, e.g.

```
shared_buffers = 16MB   # min 128KB, 16*8KB buffers
```

... PostgreSQL Buffer Manager

Shared memory

Buffer Pool Data Types

```
typedef struct buftag {
    RelFileNode rnode;      /* physical relation identifier */
    ForkNumber  forkNum;
    BlockNumber blockNum;  /* relative to start of reln */
} BufferTag;

typedef struct BufferDesc { (simplified)
    BufferTag tag;          // ID of page contained in buffer
    int       buf_id;      // buffer's index number (from 0)
    Bits32    state;       // dirty, refcount, usage
    int       freeNext;    // link in freelist chain
    ...                     // others related to concurrency
} BufferDesc;
```

Buffer Pool Functions

Buffer manager interface:

Buffer ReadBuffer(Relation r, BlockNumber n)

- ensures page *n* of file for relation `r` is loaded
- increments reference (pin) count and usage count for buffer
- returns index of loaded page in buffer pool (`Buffer` value)

BufferDesc *BufferAlloc(
            Relation r, ForkNumber f,
            BlockNumber n, bool *found)

- used by **ReadBuffer** to find a buffer for *(r,f,n)*
- if no available buffers, select buffer to be replaced

... Buffer Pool Functions

Buffer manager interface (cont):

void ReleaseBuffer(Buffer buf)

- decrement pin count on buffer
- if pin count falls to zero,
  ensures all activity on buffer is completed before returning

void MarkBufferDirty(Buffer buf)

- marks a buffer as modified
- requires that buffer is pinned and locked
- actual write is done later (e.g. when buffer replaced)

... Buffer Pool Functions

Additional buffer manager functions:

Page BufferGetPage(Buffer buf)

- finds actual data associated with buffer in pool
- returns reference to memory where data is located

BufferIsPinned(Buffer buf)

- check whether this backend holds a pin on buffer

CheckPointBuffers

- write data in checkpoint logs (for recovery)
- flush all dirty blocks in buffer pool to disk

etc. etc. etc.

Clock–sweep Replacement Strategy

PostgreSQL page replacement strategy: *clock–sweep*

- treat buffer pool as circular list of buffer slots
- `NextVictimBuffer` holds index of next possible evictee
- if this page is pinned or "popular", leave it
    - `usage_count` implements "popularity/recency" measure
    - incremented on each access to buffer (up to small limit)
    - decremented each time considered for eviction
- increment `NextVictimBuffer` and try again (wrap at end)

Exercise 6: PostgreSQL Buffer Pool

Consider an initally empty buffer pool with only 3 slots.

Show the state of the pool after each of the following:

```
Req R0, Req S0, Rel S0, Req S1, Rel S1, Req S2,
Rel S2, Rel R0, Req R1, Req S0, Rel S0, Req S1,
Rel S1, Req S2, Rel S2, Rel R1, Req R2, Req S0,
Rel S0, Req S1, Rel S1, Req S2, Rel S2, Rel R2
```
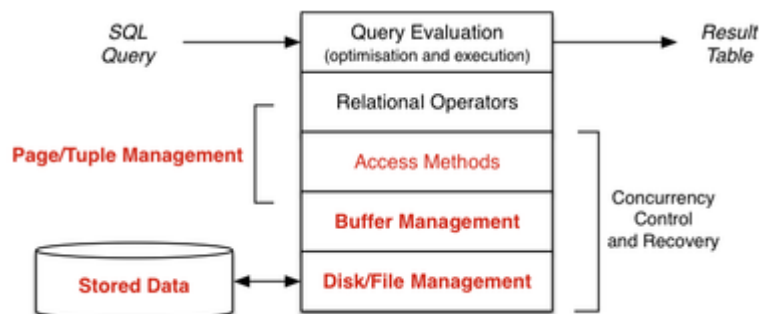
Treat `BufferDesc` entries as

`(tag, usage_count, refcount, freeNext)`

Assume `freeList` and `nextVictim` global variables.

Pages

Page/Tuple Management

Pages

Database applications view data as:

- a collection of records (tuples)
- records can be accessed via a `TupleId/RecordId/RID`
- `TupleId = (PageID + TupIndex)`

The disk and buffer manager provide the following view:

- data is a sequence of fixed–size pages (aka "blocks")
- pages can be (random) accessed via a `PageID`
- each page contains zero or more tuple values

Page format = how space/tuples are organised within a page

Page Formats

Ultimately, a `Page` is simply an array of bytes (`byte[]`).

We want to interpret/manipulate it as a collection of `Records`.

Typical operations on `Pages`:

- `request_page(pid)` ... get page via its `PageId`
- `get_record(rid)` ... get record via its `TupleId`
- `rid = insert_record(pid,rec)` ... add new record
- `update_record(rid,rec)` ... update value of record
- `delete_record(rid)` ... remove record from page

Note: `rid` contains `(PageId,TupIndex)`, so no explicit `pid` needed

... Page Formats

Factors affecting `Page` formats:

- determined by record size flexibility   (fixed, variable)
- how free space within `Page` is managed
- whether some data is stored outside `Page`
  - does `Page` have an associated overflow chain?
  - are large data values stored elsewhere? (e.g. TOAST)
 - can one tuple span multiple `Pages`?

Implementation of `Page` operations critically depends on format.

... Page Formats

For fixed–length records, use *record slots*.

- *insert*: place new record in first available slot
- *delete*: two possibilities for handling free record slots:

Exercise 7: Fixed–length Records

Give examples of table definitions

- which result in fixed–length records
- which result in variable–length records

```
create table R ( ...);
```

What are the common features of each type of table?

Exercise 8: Inserting/Deleting Fixed–length Records

For each of the following Page formats:

- compacted/packed free space
- unpacked free space (with bitmap)

Implement

- a suitable data structure to represent a `Page`
- a function to insert a new record
- a function to delete a record

Page Formats

For variable-length records, must use *slot directory*.

Possibilities for handling free-space within block:

- compacted (one region of free space)
- fragmented (distributed free space)

In practice, a combination is useful:

- normally fragmented (cheap to maintain)
- compacted when needed (e.g. record won't fit)

Important aspect of using slot directory

- location of tuple within page can change, tuple index does not change


... Page Formats

Compacted free space:

Note: "pointers" are implemented as word offsets within block.



... Page Formats

Fragmented free space:

Initial page state (compacted free space) ...

Before inserting record 7 (compacted free space) ...
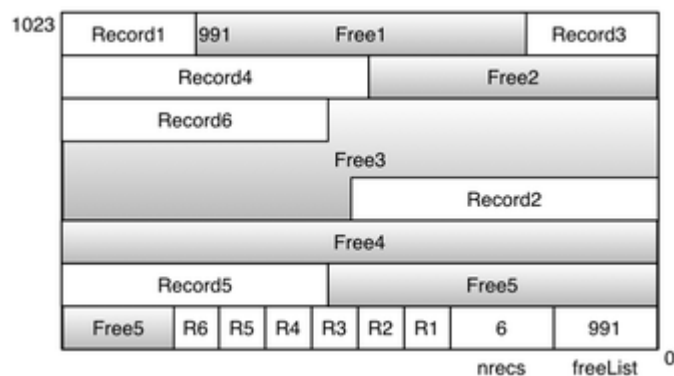
After inserting record 7 (80 bytes) ...

Initial page state (fragmented free space) ...

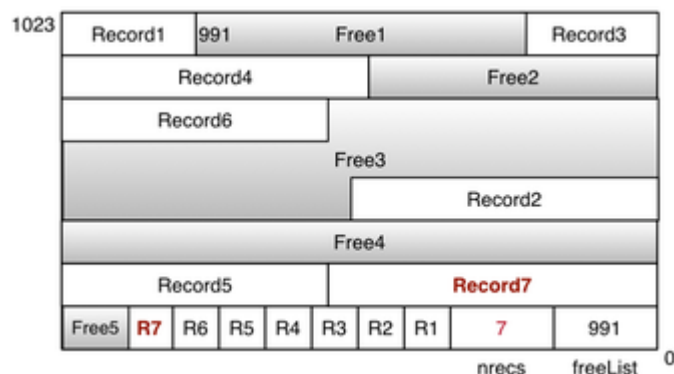Before inserting record 7 (fragmented free space) ...

After inserting record 7 (80 bytes) ...

Exercise 9: Inserting Variable–length Records

For both of the following page formats

1. variable–length records, with compacted free space
2. variable–length records, with fragmented free space

implement the `insert()` function.

Use the above page format, but also assume:

- page size is 1024 bytes
- tuples start on 4–byte boundaries
- references into page are all 8–bits (1 byte) long
- a function `recSize(r)` gives size in bytes

Storage Utilisation

How many records can fit in a page? (denoted C = capacity)

Depends on:

- page size ... typical values: 1KB, 2KB, 4KB, 8KB
- record size ... typical values: 64B, 200B, app–dependent
- page header data ... typically: 4B – 32B
- slot directory ... depends on how many records

We typically consider *average* record size (R)

Given C, $HeaderSize + C*SlotSize + C*R \leqslant PageSize$

Exercise 10: Space Utilisation

Consider the following page/record information:

- page size = 1KB = 1024 bytes = 2^{10} bytes
- records: `(a:int,b:varchar(20),c:char(10),d:int)`
- records are all aligned on 4–byte boundaries
- `c` field padded to ensure `d` starts on 4–byte boundary
- each records has 4 field–offsets at start of record (each 1 byte)
- `char(10)` field rounded up to 12–bytes to preserve alignment
- maximum size of `b` values = 20 bytes; average size = 16 bytes
- page has 32–bytes of header information, starting at byte 0
- only insertions, no deletions or updates

Calculate C = average number of records per page.

Overflows

Sometimes, it may not be possible to insert a record into a page:

1. no free–space fragment large enough
2. overall free–space is not large enough
3. the record is larger than the page
4. no more free directory slots in page

For case (1), can first try to compact free–space within the page.

If still insufficient space, we need an alternative solution ...

... Overflows

File organisation determines how cases (2)..(4) are handled.

If records may be inserted anywhere that there is free space

- cases (2) and (4) can be handled by making a new page
- case (3) requires either spanned records or "overflow file"

If file organisation determines record placement (e.g. hashed file)

- cases (2) and (4) require an "overflow page"
- case (3) requires an "overflow file"

With overflow pages, *rid* structure may need modifying *(rel,page,ovfl,rec)*

... Overflows

Overflow files for very large records and BLOBs:

Record−based handling of overflows:

We discuss overflow pages in more detail when covering Hash Files.

PostgreSQL Page Representation

Functions: **src/backend/storage/page/*.c**

Definitions: **src/include/storage/bufpage.h**

Each page is 8KB (default BLCKSZ) and contains:

- header (free space pointers, flags, xact data)
- array of (offset,length) pairs for tuples in page
- free space region (between array and tuple data)
- actual tuples themselves (inserted from end towards start)
- (optionally) region for special data (e.g. index data)

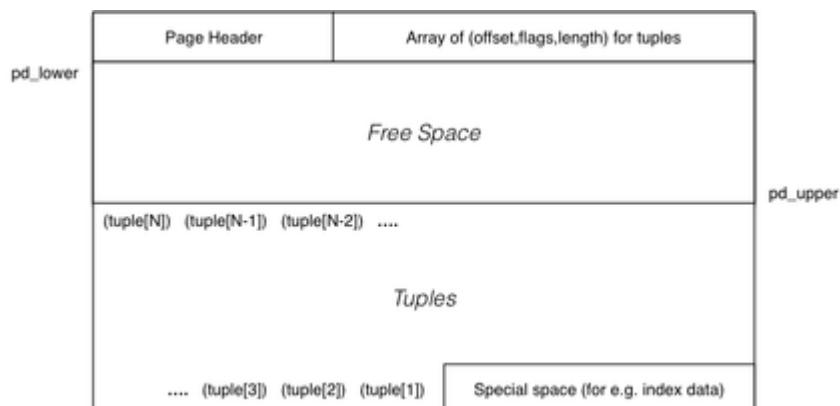Large data items are stored in separate (TOAST) files   (implicit)

Also supports ~SQL−standard BLOBs   (explicit large data items)

... PostgreSQL Page Representation

PostgreSQL page layout:

... PostgreSQL Page Representation

Page−related data types:

```
// a Page is simply a pointer to start of buffer
typedef Pointer Page;

// indexes into the tuple directory
typedef uint16  LocationIndex;

// entries in tuple directory (line pointer array)
typedef struct ItemIdData
{
   unsigned   lp_off:15,    // tuple offset from start of page
              lp_flags:2,   // unused,normal,redirect,dead
              lp_len:15;    // length of tuple (bytes)
} ItemIdData;
```

... PostgreSQL Page Representation

Page−related data types: (cont)

```
typedef struct PageHeaderData
{
    XLogRecPtr    pd_lsn;        // xact log record for last change
    uint16        pd_tli;        // xact log reference information
    uint16        pd_flags;      // flag bits (e.g. free, full, ...
    LocationIndex pd_lower;      // offset to start of free space
    LocationIndex pd_upper;      // offset to end of free space
    LocationIndex pd_special;    // offset to start of special space
    uint16        pd_pagesize_version;
    TransactionId pd_prune_xid;  // is pruning useful in data page?
    ItemIdData    pd_linp[1];    // beginning of line pointer array
} PageHeaderData;

typedef PageHeaderData *PageHeader;
```

... PostgreSQL Page Representation

Operations on `Pages`:

void PageInit(Page page, Size pageSize, ...)

- initialize a `Page` buffer to empty page
- in particular, sets `pd_lower` and `pd_upper`

OffsetNumber PageAddItem(Page page,
                          Item item, Size size, ...)

- insert one tuple (or index entry) into a `Page`
- fails if: not enough free space, too many tuples

void PageRepairFragmentation(Page page)

- compact tuple storage to give one large free space region

... PostgreSQL Page Representation

PostgreSQL has two kinds of pages:

- *heap pages* which contain tuples
- *index pages* which contain index entries

Both kinds of page have the same page layout.

One important difference:

- index entries tend be a smaller than tuples
- can typically fit more index entries per page

Exercise 11: PostgreSQL Pages

Draw diagrams of a PostgreSQL heap page

- when it is initially empty
- after three tuples have been inserted
    - with lengths of 60, 80, and 70 bytes
- after the 80 byte tuple is deleted (but before vacuuming)
- after a new 50 byte tuple is added

Show the values in the tuple header.

Assume that there is no special space in the page.

Produced: 2 Mar 2020