

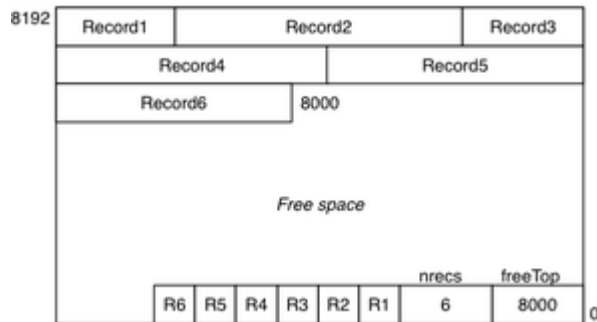
# Week 04 Lectures

## Tuples

### Tuples

2/84

Each *page* contains a collection of *tuples*



What do tuples contain? How are they structured internally?

### Records vs Tuples

3/84

A *table* is defined by a *schema*, e.g.

```
create table Employee (  
    id    integer primary key,  
    name  varchar(20) not null,  
    job   varchar(10),  
    dept  number(4) references Dept(id)  
);
```

where a schema is a collection of attributes (name,type,constraints)

Schema information (meta-data) is stored in the DB catalog

### ... Records vs Tuples

4/84

*Tuple* = collection of attribute values based on a schema, e.g.

(33357462, 'Neil Young', 'Musician', 0277)

*Record* = sequence of bytes, containing data for one tuple, e.g.

01101001	11001100	01010101	00111100	10100011	01011111	01011010
----------	----------	----------	----------	----------	----------	----------

Bytes need to be interpreted relative to schema to get tuple

### Converting Records to Tuples

5/84


A Record is an array of bytes (byte[ ])

- representing the data values from a typed Tuple

- stored on disk (persistent) or in a memory buffer

A `Tuple` is a collection of named, typed values (cf. C `struct`)

- to manipulate the values, need an "interpretable" structure
- stored in working memory, and temporary



## ... Converting Records to Tuples

6/84

Information on how to interpret bytes in a record ...

- may be contained in schema data in DBMS catalog
- may be stored in the page directory
- may be stored in the record (in a record header)
- may be stored partly in the record and partly in the schema

For variable-length records, some formatting info ...

- must be stored in the record or in the page directory
- at the least, need to know how many bytes in each value

## Operations on Records

7/84

Common operation on records ... access record via `RecordId`:

```
Record get_record(Relation rel, RecordId rid) {
    (pid,tid) = rid;
    Page buf = get_page(rel, pid);
    return get_bytes(rel, buf, tid);
}
```

Cannot use a `Record` directly; need a `Tuple`:

```
Relation rel = ... // relation schema
Record rec = get_record(rel, rid)
Tuple t = mkTuple(rel, rec)
```

Once we have a `Tuple`, we can access individual attributes/fields

## Operations on Tuples

8/84

Once we have a record, we need to interpret it as a tuple ...

```
Tuple t = mkTuple(rel, rec)
```

- convert record to tuple data structure for relation `rel`

Once we have a tuple, we want to examine its contents ...

```
Typ getTypField(Tuple t, int i)
```

- extract the `i`'th field from a `Tuple` as a value of type `Typ`


E.g. `int x = getIntField(t,1), char *s = getStrField(t,2)`

## Fixed-length Records

9/84

A possible encoding scheme for fixed-length records:

- record format (length + offsets) stored in catalog
- data values stored in fixed-size slots in data pages




Since record format is frequently used at query time, cache in memory.

## Variable-length Records

10/84

Possible encoding schemes for variable-length records:


- Prefix each field by length



- Terminate fields by delimiter



- Array of offsets



## Data Types

11/84

DBMSs typically define a fixed set of base types, e.g.

`DATE, FLOAT, INTEGER, NUMBER(n), VARCHAR(n), ...`

This determines implementation-level data types for field values:

DATE	time_t
FLOAT	float,double
INTEGER	int,long
NUMBER(n)	int[] (?)
VARCHAR(n)	char[]

PostgreSQL allows new base types to be added

A `Tuple` could be implemented as

- a list of field descriptors for a record instance  
(where a `FieldDesc` gives (offset,length,type) information)
- along with a reference to the `Record` data

```
typedef struct {
    ushort    nfields;    // number of fields/attrs
    ushort    data_off;   // offset in struct for data
    FieldDesc fields[];   // field descriptions
    Record     data;      // pointer to record in buffer
} Tuple;
```

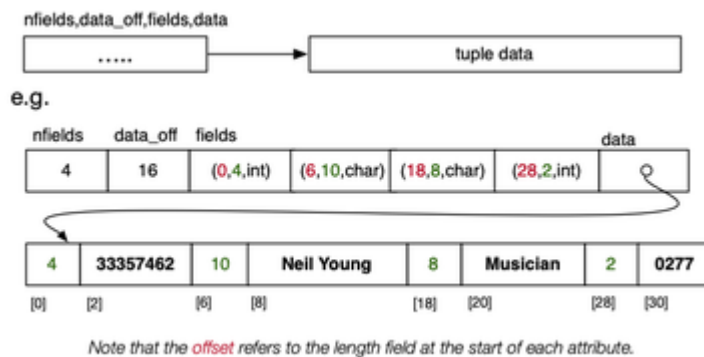
Fields are derived from relation descriptor + record instance data.

### ... Field Descriptors

13/84

Tuple data could be

- a pointer to bytes stored elsewhere in memory

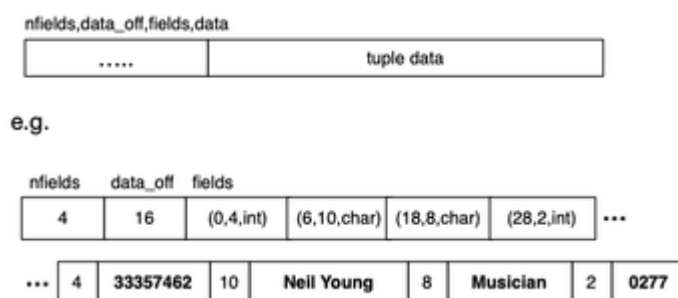


### ... Field Descriptors

14/84

Or, tuple data could be ...

- appended to `Tuple struct` (used widely in PostgreSQL)



## Exercise 1: How big is a `FieldDesc`?

15/84

`FieldDesc` = (offset,length,type), where

- offset = offset of field within record data
- length = length (in bytes) of field
- type = data type of field

If pages are 8KB in size, how many bits are needed for each?

E.g.

nfields	data_off	fields = FieldDesc[4]			
4	16	(0,4,int)	(6,10,char)	(18,8,char)	(28,2,int)

## PostgreSQL Tuples

16/84

Definitions: `include/postgres.h`, `include/access/*tup*.h`

Functions: `backend/access/common/*tup*.c` e.g.

- `HeapTuple heap_form_tuple(desc, values[], isnull[])`
- `heap_deform_tuple(tuple, desc, values[], isnull[])`

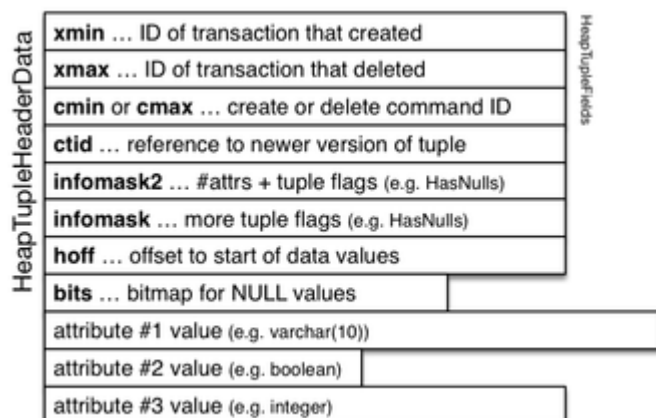
PostgreSQL implements tuples via:

- a contiguous chunk of memory
- starting with a header giving e.g. #fields, nulls
- followed by data values (as a sequence of Datum)

### ... PostgreSQL Tuples

17/84

Tuple structure:



### ... PostgreSQL Tuples

18/84

Tuple-related data types: (cont)

// TupleDesc: schema-related information for HeapTuples

```
typedef struct tupleDesc
{
    int          natts;           // # attributes in tuple
    Oid          tdtypeid;        // composite type ID for tuple type
    int32        tdtypmod;        // typmod for tuple type
    bool         tdhasoid;        // does tuple have oid attribute?
    int          tdrefcount;       // reference count (-1 if not counting)
    TupleConstr *constr;          // constraints, or NULL if none
    FormData_pg_attribute attrs[];
    // attrs[N] is a pointer to description of attribute N+1
} *TupleDesc;
```

### ... PostgreSQL Tuples

19/84

Tuple-related data types: (cont)

```
// FormData_pg_attribute:
// schema-related information for one attribute

typedef struct FormData_pg_attribute
{
    Oid      attrelid;    // OID of reln containing attr
    NameData attname;    // name of attribute
    Oid      atttypid;    // OID of attribute's data type
    int16    attlen;      // attribute length
    int32    attndims;    // # dimensions if array type
    bool     attnotnull;  // can attribute have NULL value
    .....              // and many other fields
} FormData_pg_attribute;
```

For details, see include/catalog/pg\_attribute.h

---

### ... PostgreSQL Tuples

20/84

HeapTupleData contains information about a stored tuple

```
typedef HeapTupleData *HeapTuple;

typedef struct HeapTupleData
{
    uint32      t_len;    // length of *t_data
    ItemPointerData t_self; // SelfItemPointer
    Oid          t_tableOid; // table the tuple came from
    HeapTupleHeader t_data; // -> tuple header and data
} HeapTupleData;
```

HeapTupleHeader is a pointer to a location in a buffer

---

### ... PostgreSQL Tuples

21/84

PostgreSQL stores a single block of data for tuple

- containing a tuple header, followed by data byte[ ]

```
typedef struct HeapTupleHeaderData // simplified
{
    HeapTupleFields t_heap;
    ItemPointerData t_ctid;    // TID of newer version
    uint16          t_infomask2; // #attributes + flags
    uint16          t_infomask;  // flags e.g. has_null
    uint8           t_hoff;      // sizeof header incl. t_bits
    // above is fixed size (23 bytes) for all heap tuples
    bits8           t_bits[1];   // bitmap of NULLs, var.len.
    // OID goes here if HEAP_HASOID is set in t_infomask
    // actual data follows at end of struct
} HeapTupleHeaderData;
```

---

### ... PostgreSQL Tuples

22/84

Some of the bits in t\_infomask ..

```
#define HEAP_HASNULL      0x0001
/* has null attribute(s) */
#define HEAP_HASVARWIDTH  0x0002
/* has variable-width attribute(s) */
#define HEAP_HASEXTERNAL  0x0004
/* has external stored attribute(s) */
#define HEAP_HASOID_OLD   0x0008
/* has an object-id field */
```

## ... PostgreSQL Tuples

23/84

Tuple-related data types: (cont)

```
typedef struct HeapTupleFields // simplified
{
    TransactionId t_xmin; // inserting xact ID
    TransactionId t_xmax; // deleting or locking xact ID
    union {
        CommandId t_cid; // inserting or deleting command ID
        TransactionId t_xvac; // old-style VACUUM FULL xact ID
    } t_field3;
} HeapTupleFields;
```

Note that not all system fields from stored tuple appear

- oid is stored after the tuple header, if used
- both xmin/xmax are stored, but only one of cmin/cmax

---

## PostgreSQL Attribute Values

24/84

Values of attributes in PostgreSQL tuples are packaged as Datums

```
// representation of a data value
typedef uintptr_t Datum;
```

The actual data value:

- may be stored in the Datum (e.g. int)
- may have a header with length (for varlen attributes)
- may be stored in a TOAST file (if large value)

---

## ... PostgreSQL Attribute Values

25/84

Attribute values can be extracted as Datum from HeapTuples

```
Datum heap_getattr(
    HeapTuple tup, // tuple (in memory)
    int attnum, // which attribute
    TupleDesc tupDesc, // field descriptors
    bool *isnull // flag to record NULL
)
```

isnull is set to true if value of field is NULL

attnum can be negative ... to access system attributes (e.g. OID)

For details, see include/access/htup\_details.h

---

## ... PostgreSQL Attribute Values

26/84

Values of Datum objects can be manipulated via e.g.

```
// DatumGetBool:
// Returns boolean value of a Datum.

#define DatumGetBool(X) ((bool) ((X) != 0))

// BoolGetDatum:
```

```
// Returns Datum representation for a boolean.
```

```
#define BoolGetDatum(X) ((Datum) ((X) ? 1 : 0))
```

For details, see include/postgres.h

---

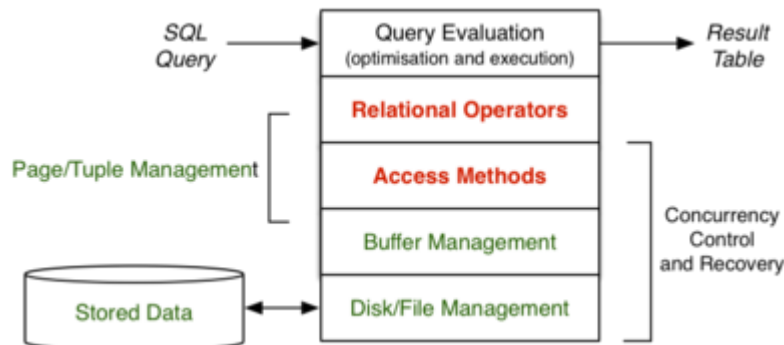
## Implementing Relational Operations

---

### DBMS Architecture (revisited)

28/84

Implementation of relational operations in DBMS:



---

### Relational Operations

29/84

DBMS core = relational engine, with implementations of

- selection, projection, join, set operations
- scanning, sorting, grouping, aggregation, ...

In this part of the course:

- examine methods for implementing each operation
- develop cost models for each implementation
- characterise when each method is most effective

Terminology reminder:

- tuple = collection of data values under some schema  $\equiv$  record
- page = block = collection of tuples + management data = i/o unit
- relation = table  $\equiv$  file = collection of tuples

---

### ... Relational Operations

30/84

Two "dimensions of variation":

- which relational operation (e.g. Sel, Proj, Join, Sort, ...)
- which access-method (e.g. file struct: heap, indexed, hashed, ...)

Each *query method* involves an operator and a file structure:

- e.g. primary-key selection on hashed file
- e.g. primary-key selection on indexed file
- e.g. join on ordered heap files (sort-merge join)
- e.g. join on hashed files (hash join)
- e.g. two-dimensional range query on R-tree indexed file



As well as query costs, consider update costs (insert/delete).

---

## ... Relational Operations

31/84

SQL vs DBMS engine

- **select ... from R where C**
  - find relevant tuples (satisfying C) in file(s) of R
- **insert into R values(...)**
  - place new tuple in some page of a file of R
- **delete from R where C**
  - find relevant tuples and "remove" from file(s) of R
- **update R set ... where C**
  - find relevant tuples in file(s) of R and "change" them

---

## Cost Models

---

### Cost Models

33/84

An important aspect of this course is

- analysis of cost of various query methods

*Cost* can be measured in terms of

- *Time Cost*: total time taken to execute method, or
- *Page Cost*: number of pages read and/or written

Primary assumptions in our cost models:

- memory (RAM) is "small", fast, byte-at-a-time
- disk storage is very large, slow, page-at-a-time

---

### ... Cost Models

34/84

Since *time cost* is affected by many factors

- speed of i/o devices (fast/slow disk, SSD)
- load on machine

we do not consider time cost in our analyses.

For comparing methods, *page cost* is better

- identifies workload imposed by method
- BUT is clearly affected by buffering

Estimating costs with multiple concurrent ops *and* buffering is difficult!!

Additional assumption: every page request leads to some i/o

---

### ... Cost Models

35/84

In developing cost models, we also assume:

- a relation is a set of  $r$  tuples, with average size  $R$  bytes
- the tuples are stored in  $b$  data pages on disk
- each page has size  $B$  bytes and contains up to  $c$  tuples
- the tuples which answer query  $q$  are contained in  $b_q$  pages
- data is transferred disk↔memory in whole pages
- cost of disk↔memory transfer  $T_{r/w}$  is very high



## ... Cost Models

36/84

Our cost models are "rough" (based on assumptions)

But do give an  $O(x)$  feel for how expensive operations are.

Example "rough" estimation: how many piano tuners in Sydney?

- Sydney has  $\approx 4\,000\,000$  people
- Average household size  $\approx 3 \therefore 1\,300\,000$  households
- Let's say that 1 in 10 households owns a piano
- Therefore there are  $\approx 130\,000$  pianos
- Say people get their piano tuned every 2 years (on average)
- Say a tuner can do 2/day, 250 working-days/year
- Therefore 1 tuner can do 500 pianos per year
- Therefore Sydney would need  $\approx 130000/2/500 = 130$  tuners

Actual number of tuners in Yellow Pages = 120

Example borrowed from Alan Fekete at Sydney University.

## Query Types

37/84

Type	SQL	RelAlg	a.k.a.
Scan	<code>select * from R</code>	$R$	–
Proj	<code>select x,y from R</code>	$Proj[x,y]R$	–
Sort	<code>select * from R order by x</code>	$Sort[x]R$	<i>ord</i>
$Sel_1$	<code>select * from R where id = k</code>	$Sel[id=k]R$	<i>one</i>
$Sel_n$	<code>select * from R where a = k</code>	$Sel[a=k]R$	–
$Join_1$	<code>select * from R,S where R.id = S.r</code>	$R Join[id=r] S$	–


Different query classes exhibit different query processing behaviours.

## Example File Structures

38/84

When describing file structures

- use a large box to represent a *page*
- use either a small box or  $tup_i$  (or  $rec_i$ ) to represent a *tuple*
- sometimes refer to tuples via their *key*
  - mostly, *key* corresponds to the notion of "primary key"
  - sometimes, *key* means "search key" in selection condition



## ... Example File Structures

39/84

Consider three simple file structures:

- *heap file* ... tuples added to any page which has space
- *sorted file* ... tuples arranged in file in key order
- *hash file* ... tuples placed in pages using hash function

All files are composed of  $b$  primary blocks/pages



Some records in each page may be marked as "deleted".

## Exercise 2: Operation Costs

40/84

For each of the following file structures

- heap file, sorted file, hash file

Determine #page-reads + #page-writes for insert and delete

You can assume the existence of a file header containing

- values for  $r, R, b, B, c$
- index of first page with free space (and a free list)

Assume also

- each page contains a header and directory as well as tuples
- no buffering (worst case scenario)

## Scanning

### Scanning

42/84

Consider the query:

```
select * from Rel;
```

Operational view:

```
for each page P in file of relation Rel {  
  for each tuple t in page P {  
    add tuple t to result set  
  }  
}
```


Cost: read every data page once

$Time\ Cost = b \cdot T_P$ ,  $Page\ Cost = b$

### ... Scanning

43/84

Scan implementation when file has overflow pages, e.g.



### ... Scanning

44/84

In this case, the implementation changes to:

```
for each page P in data file of relation Rel {  
  for each tuple t in page P {  
    add tuple t to result set  
  }  
  for each overflow page V of page P {  
    for each tuple t in page V {  
      add tuple t to result set  
    }  
  }  
}
```

Cost: read each data page and each overflow page once

$Cost = b + b_{OV}$

where  $b_{OV}$  = total number of overflow pages


## Selection via Scanning

45/84

Consider a *one* query like:

```
select * from Employee where id = 762288;
```

In an unordered file, search for matching tuple requires:



Guaranteed at most one answer; but could be in any page.

---

## ... Selection via Scanning

46/84

Overview of scan process:

```
for each page P in relation Employee {
  for each tuple t in page P {
    if (t.id == 762288) return t
  }
}
```

Cost analysis for *one* searching in unordered file

- best case: read one page, find tuple
- worst case: read all  $b$  pages, find in last (or don't find)
- average case: read half of the pages ( $b/2$ )


Page Costs:  $Cost_{avg} = b/2$     $Cost_{min} = 1$     $Cost_{max} = b$

---

## Exercise 3: Cost of Search in Hashed File

47/84

Consider the hashed file structure  $b = 10$ ,  $c = 4$ ,  $h(k) = k \% 10$



Describe how the following queries

```
select * from R where k = 51;
select * from R where k > 50;
```

might be solved in a file structure like the above ( $h(k) = k \% b$ ).

Estimate the minimum and maximum cost (as #pages read)

---

## Iterators

48/84

Access methods typically involve *iterators*, e.g.

**Scan s = start\_scan(Relation r, ...)**

- commence a scan of relation  $r$
- Scan may include condition to implement WHERE-clause
- Scan holds data on progress through file (e.g. current page)

**Tuple next\_tuple(Scan s)**

- return Tuple immediately following last accessed one
- returns NULL if no more Tuples left in the relation

---

49/84

## Example Query

Example: simple scan of a table ...

```
select name from Employee
```

implemented as:

```
DB db = openDatabase("myDB");
Relation r = openRelation(db, "Employee", READ);
Scan s = start_scan(r);
Tuple t; // current tuple
while ((t = next_tuple(s)) != NULL)
{
    char *name = getStrField(t, 2);
    printf("%s\n", name);
}
```

---

## Exercise 4: Implement next\_tuple()

50/84

Consider the following possible **Scan** data structure

```
typedef struct {
    Relation rel;
    Page *curPage; // Page buffer
    int curPID; // current pid
    int curTID; // current tid
} ScanData;
```

Assume tuples are indexed 0..nTuples(p)

Assume pages are indexed 0..nPages(rel)

Implement the **Tuple next\_tuple(Scan)** function

P.S. What's in a Relation object?

---

## Relation Copying

51/84

Consider an SQL statement like:

```
create table T as (select * from S);
```

Effectively, copies data from one table to a new table.

Process:

```
make empty relation T
s = start scan of S
while (t = next_tuple(s)) {
    insert tuple t into relation T
}
```


---

### ... Relation Copying

52/84

Possible that T is smaller than S

- may be unused free space in S where tuples were removed
- if T is built by simple append, will be compact



## ... Relation Copying

53/84

In terms of existing relation/page/tuple operations:

```

Relation in;          // relation handle (incl. files)
Relation out;         // relation handle (incl. files)
int ipid,opid,tid;    // page and record indexes
Record rec;           // current record (tuple)
Page ibuf,obuf;       // input/output file buffers

in = openRelation("S", READ);
out = openRelation("T", NEW|WRITE);
clear(obuf); opid = 0;
for (ipid = 0; ipid < nPages(in); ipid++) {
    ibuf = get_page(in, ipid);
    for (tid = 0; tid < nTuples(ibuf); tid++) {
        rec = get_record(ibuf, tid);
        if (!hasSpace(obuf,rec)) {
            put_page(out, opid++, obuf);
            clear(obuf);
        }
        insert_record(obuf,rec);
    }
}
if (nTuples(obuf) > 0) put_page(out, opid, obuf);

```

## Exercise 5: Cost of Relation Copy

54/84

Analyse cost for relation copying:

1. if both input and output are heap files
2. if input is sorted and output is heap file
3. if input is heap file and output is sorted

Assume ...

- $r$  records in input file,  $c$  records/page
- $b_{in}$  = number of pages in input file
- some pages in input file are *not* full
- all pages in output file are full (except the last)

Give cost in terms of #pages read + #pages written

## Scanning in PostgreSQL

55/84

Scanning defined in: [backend/access/heap/heapam.c](#)

Implements iterator data/operations:

- **HeapScanDesc** ... struct containing iteration state
- **scan = heap\_beginscan(rel,...,nkeys,keys)**
- **tup = heap\_getnext(scan, direction)**
- **heap\_endscan(scan)** ... frees up scan struct
- **res = HeapKeyTest(tuple,...,nkeys,keys)**  
... performs ScanKeys tests on tuple ... is it a result tuple?

```
typedef HeapScanDescData *HeapScanDesc;

typedef struct HeapScanDescData
{
    // scan parameters
    Relation      rs_rd;           // heap relation descriptor
    Snapshot      rs_snapshot;     // snapshot ... tuple visibility
    int           rs_nkeys;        // number of scan keys
    ScanKey       rs_key;          // array of scan key descriptors
    ...
    // state set up at initscan time
    PageNumber    rs_npages;       // number of pages to scan
    PageNumber    rs_startpage;    // page # to start at
    ...
    // scan current state, initially set to invalid
    HeapTupleData rs_ctup;         // current tuple in scan
    PageNumber    rs_cpage;        // current page # in scan
    Buffer         rs_cbuf;         // current buffer in scan
    ...
} HeapScanDescData;
```

---

## Scanning in other File Structures

57/84

Above examples are for *heap* files

- simple, unordered, maybe indexed, no hashing

Other access file structures in PostgreSQL:

- **btree**, **hash**, **gist**, **gin**
  - each implements:
    - startscan, getnext, endscan
    - insert, delete (update=delete+insert)
    - other file-specific operators
- 

## Sorting

### The Sort Operation

59/84

Sorting is explicit in queries only in the `order by` clause

```
select * from Students order by name;
```

Sorting is used internally in other operations:

- eliminating duplicate tuples for projection
- ordering files to enhance select efficiency
- implementing various styles of join
- forming tuple groups in `group by`

Sort methods such as quicksort are designed for in-memory data.

For large data on disks, need external sorts such as *merge sort*.

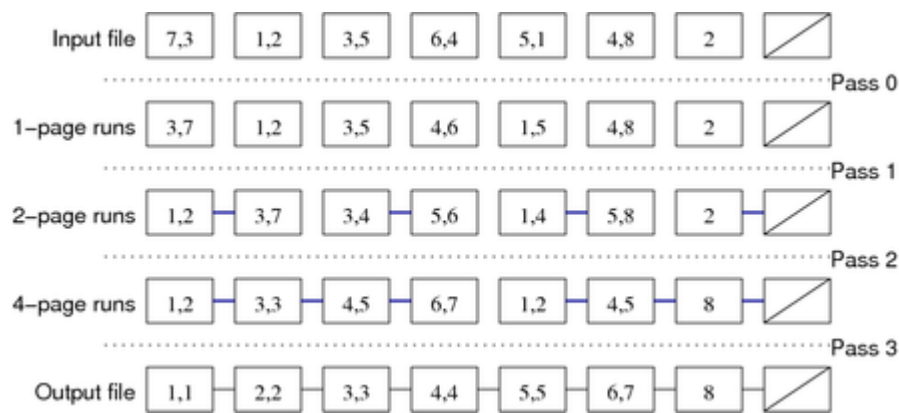
---

### Two-way Merge Sort

60/84

Example:






## ... Two-way Merge Sort

61/84

Requires three in-memory buffers:



Assumption: cost of Merge operation on two in-memory buffers  $\approx 0$ .

## Comparison for Sorting

62/84

Above assumes that we have a function to compare tuples.

Needs to understand ordering on different data types.

Need a function `tupCompare(r1,r2,f)` (cf. C's `strcmp`)

```
int tupCompare(r1,r2,f)
{
    if (r1.f < r2.f) return -1;
    if (r1.f > r2.f) return 1;
    return 0;
}
```

Assume `=`, `<`, `>` are available for all attribute types.

## ... Comparison for Sorting

63/84

In reality, need to sort on multiple attributes and ASC/DESC, e.g.

```
-- example multi-attribute sort
select * from Students
order by age desc, year_enrolled
```

Sketch of multi-attribute sorting function

```
int tupCompare(r1,r2,criteria)
{
    foreach (f,ord) in criteria {
        if (ord == ASC) {
            if (r1.f < r2.f) return -1;
            if (r1.f > r2.f) return 1;
        }
        else {
            if (r1.f > r2.f) return -1;
        }
    }
}
```

```

    if (r1.f < r2.f) return 1;
  }
}
return 0;
}

```

## Cost of Two-way Merge Sort

64/84

For a file containing  $b$  data pages:

- require  $\text{ceil}(\log_2 b)$  passes to sort,
- each pass requires  $b$  page reads,  $b$  page writes

Gives total cost:  $2 \cdot b \cdot \text{ceil}(\log_2 b)$

Example: Relation with  $r=10^5$  and  $c=50 \Rightarrow b=2000$  pages.


Number of passes for sort:  $\text{ceil}(\log_2 2000) = 11$

Reads/writes entire file 11 times! Can we do better?

## n-Way Merge Sort

65/84

Initial pass uses:  $B$  total buffers




Reads  $B$  pages at a time, sorts in memory, writes out in order

### ... n-Way Merge Sort

66/84

Merge passes use:  $n$  input buffers,  $1$  output buffer



### ... n-Way Merge Sort

67/84

Method:

```

// Produce B-page-long runs
for each group of B pages in Rel {
    read B pages into memory buffers
}

```

```

    sort group in memory
    write B pages out to Temp
}
// Merge runs until everything sorted
numberOfRuns = ⌈b/B⌉
while (numberOfRuns > 1) {
    // n-way merge, where n=B-1
    for each group of n runs in Temp {
        merge into a single run via input buffers
        write run to newTemp via output buffer
    }
    numberOfRuns = ⌈numberOfRuns/n⌉
    Temp = newTemp // swap input/output files
}

```

---

## Cost of n-Way Merge Sort

68/84

Consider file where  $b = 4096$ ,  $B = 16$  total buffers:

- pass 0 produces  $256 \times 16$ -page sorted runs
- pass 1
  - performs 15-way merge of groups of 16-page sorted runs
  - produces  $18 \times 240$ -page sorted runs (17 full runs, 1 short run)
- pass 2
  - performs 15-way merge of groups of 240-page sorted runs
  - produces  $2 \times 3600$ -page sorted runs (1 full run, 1 short run)
- pass 3
  - performs 15-way merge of groups of 3600-page sorted runs
  - produces  $1 \times 4096$ -page sorted runs

(cf. two-way merge sort which needs 11 passes)

---

### ... Cost of n-Way Merge Sort

69/84

Generalising from previous example ...

For  $b$  data pages and  $B$  buffers

- first pass: read/writes  $b$  pages, gives  $b_0 = \lceil b/B \rceil$  runs
- then need  $\lceil \log_n b_0 \rceil$  passes until sorted, where  $n = B-1$
- each pass reads and writes  $b$  pages (i.e.  $2b$  page accesses)

$Cost = 2b \cdot (1 + \lceil \log_n b_0 \rceil)$ , where  $b_0$  and  $n$  are defined above

---

## Exercise 6: Cost of n-Way Merge Sort

70/84

How many reads+writes to sort the following:

- $r = 1048576$  tuples ( $2^{20}$ )
- $R = 62$  bytes per tuple (fixed-size)
- $B = 4096$  bytes per page
- $H = 96$  bytes of header data per page
- $D = 1$  presence bit per tuple in page directory
- all pages are full

Consider for the cases:

- 9 total buffers, 8 input buffers, 1 output buffer
- 33 total buffers, 32 input buffers, 1 output buffer
- 257 total buffers, 256 input buffers, 1 output buffer

---

## Sorting in PostgreSQL

71/84

Sort uses a merge-sort (from Knuth) similar to above:

- [backend/utls/sort/tuplesort.c](#)
- [include/utls/sortsupport.h](#)

Tuples are mapped to **SortTuple** structs for sorting:

- containing pointer to tuple and sort key
- no need to reference actual Tuples during sort
- unless multiple attributes used in sort

If all data fits into memory, sort using **qsort()**.

If memory fills while reading, form "runs" and do disk-based sort.

---

## ... Sorting in PostgreSQL

72/84

Disk-based sort has phases:

- divide input into sorted runs using HeapSort
- merge using  $N$  buffers, one output buffer
- $N$  = as many buffers as `workMem` allows

Described in terms of "tapes" ("tape"  $\approx$  sorted run)

Implementation of "tapes": [backend/utls/sort/logtape.c](#)

---

## ... Sorting in PostgreSQL

73/84

Sorting comparison operators are obtained via catalog (in `Type.o`):

```
// gets pointer to function via pg_operator
struct Tuplesortstate { ... SortTupleComparator ... };

// returns negative, zero, positive
ApplySortComparator(Datum datum1, bool isnull1,
                    Datum datum2, bool isnull2,
                    SortSupport sort_helper);
```

Flags in `SortSupport` indicate: ascending/descending, nulls-first/last.

`ApplySortComparator()` is PostgreSQL's version of `tupCompare()`

---

# Implementing Projection

---

## The Projection Operation

75/84

Consider the query:

```
select distinct name,age from Employee;
```

If the `Employee` relation has four tuples such as:

```
(94002, John, Sales, Manager, 32)
(95212, Jane, Admin, Manager, 39)
(96341, John, Admin, Secretary, 32)
(91234, Jane, Admin, Secretary, 21)
```

then the result of the projection is:

```
(Jane, 21) (Jane, 39) (John, 32)
```

Note that duplicate tuples (e.g. `(John, 32)`) are eliminated.

---

## ... The Projection Operation

76/84

The projection operation needs to:

1. scan the entire relation as input
  - already seen how to do scanning

2. remove unwanted attributes in output tuples
  - implementation depends on tuple internal structure
  - essentially, make a new tuple with fewer attributes and where the values may be computed from existing attributes
3. eliminate any duplicates produced (if distinct)
  - two approaches: sorting or hashing

## Sort-based Projection

77/84

Requires a temporary file/relation (Temp)

```
for each tuple T in Rel {
    T' = mkTuple([attrs],T)
    write T' to Temp
}
```

```
sort Temp on [attrs]
```

```
for each tuple T in Temp {
    if (T == Prev) continue
    write T to Result
    Prev = T
}
```

## Exercise 7: Cost of Sort-based Projection

78/84

Consider a table  $R(x,y,z)$  with tuples:

```
Page 0:  (1,1,'a')  (11,2,'a')  (3,3,'c')
Page 1:  (13,5,'c') (2,6,'b')  (9,4,'a')
Page 2:  (6,2,'a')  (17,7,'a') (7,3,'b')
Page 3:  (14,6,'a') (8,4,'c')  (5,2,'b')
Page 4:  (10,1,'b') (15,5,'b') (12,6,'b')
Page 5:  (4,2,'a')  (16,9,'c') (18,8,'c')
```

SQL: create T as (select distinct y from R)

Assuming:

- 3 memory buffers, 2 for input, one for output
- pages/buffers hold 3 R tuples (i.e.  $c_R=3$ ), 6 T tuples (i.e.  $c_T=6$ )

Show how sort-based projection would execute this statement.

## Cost of Sort-based Projection

79/84

The costs involved are (assuming  $B=n+1$  buffers for sort):


- scanning original relation Rel:  $b_R$  (with  $c_R$ )
- writing Temp relation:  $b_T$  (smaller tuples,  $c_T > c_R$ , sorted)
- sorting Temp relation:  $2 \cdot b_T \cdot \text{ceil}(\log_n b_0)$  where  $b_0 = \text{ceil}(b_T/B)$
- scanning Temp, removing duplicates:  $b_T$
- writing the result relation:  $b_{Out}$  (maybe less tuples)

Cost = sum of above =  $b_R + b_T + 2 \cdot b_T \cdot \text{ceil}(\log_n b_0) + b_T + b_{Out}$

## Hash-based Projection

80/84


Partitioning phase:



### ... Hash-based Projection

81/84

Duplicate elimination phase:



### ... Hash-based Projection

82/84

Algorithm for both phases:

```

for each tuple T in relation Rel {
  T' = mkTuple([attrs],T)
  H = h1(T', n)
  B = buffer for partition[H]
  if (B full) write and clear B
  insert T' into B
}
for each partition P in 0..n-1 {
  for each tuple T in partition P {
    H = h2(T, n)
    B = buffer for hash value H
    if (T not in B) insert T into B
    // assumes B never gets full
  }
  write and clear all buffers
}

```

## Exercise 8: Cost of Hash-based Projection

83/84

Consider a table  $R(x,y,z)$  with tuples:

Page 0: (1,1,'a') (11,2,'a') (3,3,'c')  
 Page 1: (13,5,'c') (2,6,'b') (9,4,'a')  
 Page 2: (6,2,'a') (17,7,'a') (7,3,'b')  
 Page 3: (14,6,'a') (8,4,'c') (5,2,'b')  
 Page 4: (10,1,'b') (15,5,'b') (12,6,'b')  
 Page 5: (4,2,'a') (16,9,'c') (18,8,'c')

-- and then the same tuples repeated for pages 6-11

SQL: `create T as (select distinct y from R)`

Assuming:

- 4 memory buffers, one for input, 3 for partitioning
- pages/buffers hold 3 R tuples (i.e.  $c_R=3$ ), 4 T tuples (i.e.  $c_T=4$ )
- hash functions:  $h_1(x) = x\%3$ ,  $h_2(x) = (x\%4)\%3$

Show how hash-based projection would execute this statement.

---

## Cost of Hash-based Projection

84/84

The total cost is the sum of the following:

- scanning original relation R:  $b_R$
- writing partitions:  $b_P$  ( $b_R$  vs  $b_P$ ?)
- re-reading partitions:  $b_P$
- writing the result relation:  $b_{Out}$

$$\text{Cost} = b_R + 2b_P + b_{Out}$$

To ensure that  $n$  is larger than the largest partition ...

- use hash functions ( $h_1, h_2$ ) with uniform spread
- allocate at least  $\sqrt{b_R} + 1$  buffers
- if insufficient buffers, significant re-reading overhead

---

Produced: 10 Mar 2020